

Problem Set 11

Ian McGroarty

19APRIL2020

Problem 1:

We can see from the histograms below that the pair bootstrap has a longer tail. In general it seems to be a bit more skewed. Though both are centered around the same mean.

```
#### Set up data ####
R = c(68, 77, 299, 220, 142, 287, 276, 115, 64, 206, 222, 205, 233, 228, 188, 132, 285, 188, 224
, 121, 311, 166, 248, 161, 226, 67, 201, 267, 121, 301, 244, 222, 195, 203, 210, 275, 286, 275,
304, 214)

S = c(56, 62, 445, 279, 138, 428, 319, 102, 51, 289, 351, 282, 310, 266, 256, 144, 447, 186, 389
, 113, 412, 176, 313, 162, 368, 54, 214, 429, 115, 407, 265, 301, 234, 229, 270, 478, 419, 490,
430, 235)

salmon <- as.data.frame(R)
salmon$S <- S
salmon$R2 <- 1/R
salmon$S2 <- 1/S

#### First regression ####
origRgrsn <- lm(formula = R2 ~ S2, data=salmon)


#### Bootstrap Bill Turner ####
# Set Parameters
B <- 500
pairBootstrap = rep(0,B)
residBootstrap = rep(0,B)

for (j in 1:B) {
  # bootstrap covariates and response
  ## Get Sample Data - sample by row number
  resampledObs <- base::sample(seq(1:length(salmon[,1])),length(salmon[,1]), replace = T
RUE)
  newDataSet = salmon[resampledObs,]
  ## Run regression and store coefficient
  newRgrsn = lm(formula = R2 ~ S2, data=newDataSet)
  pairBootstrap[j] = newRgrsn$coefficients[2]


  # bootstrap residuals
  ## The the residual for each observation.
  resampledResids <- origRgrsn$residuals[resampledObs]
  ## Get prediction
  salmon$alteredRspns = resampledResids + predict(origRgrsn)
  newResidRgrsn = lm(alteredRspns ~ S2,data=salmon)
  residBootstrap[j] = newResidRgrsn$coefficients[2]
}

#plot(S,R)
par(mfrow = c(2,1))
hist(pairBootstrap, freq=FALSE)
hist(residBootstrap)
```

Histogram of pairBootstrap



Histogram of residBootstrap



Problem 2

Generate 100 samples X_1, \dots, X_{20} from a normal population $N(\theta, 1)$ with $\theta = 1$.

(a) For each sample compute the bootstrap and jackknife estimate for variance for $\hat{\theta} = \bar{X}$ and compute the mean and standard deviation of these variance estimates over the 100

samples.

```

B <- 1000
theta <- c()

## Build the functions
bootstrap.var <- function(DATA,B){
  for (i in 1:B){
    # Make sure the data is in vector form
    vctrzdData <- as.vector(DATA)
    # Sample of size n with replacement from data
    btstrpSmpl = base::sample(vctrzdData,length(vctrzdData), replace = TRUE)
    # Sort
    btstrpSmpl.ordered <- sort(btstrpSmpl)
    # Take the mean of the middle 4 numbers
    # theta[i] <- mean(btstrpSmpl.ordered[3:6])
    theta[i] <- mean(btstrpSmpl.ordered)
  }
  # Take the mean of the B thetas
  theta.mean <- mean(theta)
  # Take the variance of the B thetas (var(theta) also works)
  theta.var <- (1/(B-1))*(sum((theta-theta.mean)^2))
  # Return
  return(c(theta.mean,theta.var))
}

jackknife.var <- function(DATA){
  # Clear theta
  theta <- c()
  # Get the Length of DATA
  n <- length(DATA)
  # Get the theta_j values for the jackknife samples.
  for (j in 1:n){
    # Generate the sample excluding the ith value of X1
    btstrpSmpl <- DATA[-j]
    # Take the mean of the sample.
    theta[j] <- mean(btstrpSmpl)
  }
  # Compute theta(.)
  theta.dot <- mean(theta)
  # Jackknife variance
  dif <- (theta - theta.dot)^2
  var.jackknife <- sqrt(((n-1)/n)*sum(dif))
  # Return
  return(c(theta.dot,var.jackknife))
  # print(paste0("Jackknife: Mean = ",theta.dot," Variance = ", var.jackknife))
}

## Run
jack.mean <- c()
jack.var <- c()
boot.mean <- c()
boot.var <- c()
for ( i in 1:100){
  # Generate X1

```

```

X1 <- rnorm(n=20,mean=1,sd=1)
# Bootstrap
boot <- bootstrap.var(X1,B)
  boot.mean[i] <- boot[1]
  boot.var[i] <- boot[2]
# Jackknife
jack <- jackknife.var(X1)
  jack.mean[i] <- jack[1]
  jack.var[i] <- jack[2]
}

print(paste0("Bootstrap: B = ",B, " Mean = ",mean(boot.mean), " mean(Variance) = ", mean(boot.va
r), " St. Dev.(variance) = ",sd(boot.var)))

```

```
## [1] "Bootstrap: B = 1000 Mean = 0.994970100498525 mean(Variance) = 0.047877760262376 St. Dev.
(variance) = 0.0170191611810896"
```

```
print(paste0("Jackknife: Mean = ",mean(jack.mean)," mean(Variance) = ", mean(jack.var)," St. De
v.(variance) = ",sd(jack.var)))
```

```
## [1] "Jackknife: Mean = 0.994536074257312 mean(Variance) = 0.222451568199411 St. Dev.(varianc
e) = 0.0409148554100715"
```

(b) Repeat (a) for $\hat{\theta} = \bar{X}^2$

```

B <- 1000
theta <- c()

bootstrap.var <- function(DATA,B){
  for (i in 1:B){
    # Make sure the data is in vector form
    vctrzdData <- as.vector(DATA)
    # Sample of size n with replacement from data
    btstrpSmpl = base::sample(vctrzdData,length(vctrzdData), replace = TRUE)
    # Sort
    btstrpSmpl.ordered <- sort(btstrpSmpl)
    # Take the mean of the middle 4 numbers
    # theta[i] <- mean(btstrpSmpl.ordered[3:6])
    theta[i] <- mean(btstrpSmpl.ordered)
  }
  # Take the mean of the B thetas
  theta.mean <- (mean(theta))^2
  # Take the variance of the B thetas (var(theta) also works)
  theta.var <- (1/(B-1))*(sum((theta-theta.mean)^2))
  # Return
  return(c(theta.mean,theta.var))
}

jackknife.var <- function(DATA){
  # Clear theta
  theta <- c()
  # Get the Length of DATA
  n <- length(DATA)
  # Get the theta_j values for the jackknife samples.
  for (j in 1:n){
    # Generate the sample excluding the ith value of X1
    btstrpSmpl <- DATA[-j]
    # Take the mean of the sample.
    theta[j] <- (mean(btstrpSmpl))^2
  }
  # Compute theta(.)
  theta.dot <- mean(theta)
  # Jackknife variance
  dif <- (theta - theta.dot)^2
  var.jackknife <- sqrt(((n-1)/n)*sum(dif))
  # Return
  return(c(theta.dot,var.jackknife))
  # print(paste0("Jackknife: Mean = ",theta.dot," Variance = ", var.jackknife))
}

jack.mean <- c()
jack.var <- c()
boot.mean <- c()
boot.var <- c()
for ( i in 1:100){
  # Generate X1
  X1 <- rnorm(n=20,mean=1,sd=1)
  # Bootstrap

```

```

boot <- bootstrap.var(X1,B)
  boot.mean[i] <- boot[1]
  boot.var[i] <- boot[2]
# Jackknife
jack <- jackknife.var(X1)
  jack.mean[i] <- jack[1]
  jack.var[i] <- jack[2]
}

print(paste0("Bootstrap: B = ",B, " Mean = ",mean(boot.mean), " mean(Variance) = ", mean(boot.va
r), " St. Dev.(variance) = ",sd(boot.var)))

```

```
## [1] "Bootstrap: B = 1000 Mean = 1.1038158182034 mean(Variance) = 0.110581802674557 St. Dev.(v
ariance) = 0.126440650553676"
```

```
print(paste0("Jackknife: Mean = ",mean(jack.mean)," mean(Variance) = ", mean(jack.var)," St. De
v.(variance) = ",sd(jack.var)))

```

```
## [1] "Jackknife: Mean = 1.10764142696516 mean(Variance) = 0.449888186545517 St. Dev.(variance)
= 0.114055888362321"
```

Problem 3

Find the number of bootstraps necessary to accurately calculate a 95% confidence interval in the difference of two population means. That is, simulate n values of X , where $X_i \sim N(\mu_1, \sigma_1^2)$ and m values of Y , where $Y_j \sim N(\mu_2, \sigma_2^2)$ and calculate a 95% confidence interval in $\mu_1 - \mu_2$. Now calculate a 95% confidence interval for $\mu_1 - \mu_2$ using the bootstrap techniques learned in this module. How large does B have to be until you accurately and dependably reproduce the confidence interval you originally calculated?

```
# Set parameters
B <- 200
n <- 200
mu1 <- 1
sig1 <- 1
m <- 20
mu2 <- 1
sig2 <- 2

# Generate the true samples
X <- rnorm(n,mean=mu1,sd=sig1)
Y <- rnorm(m,mean=mu2,sd=sig2)

# Theoretical Confidence interval
# Sample variance (Larsen & Marx pg 475)
sample.var <- function(DATA){
  (length(DATA)*(sum(DATA^2)) - (sum(DATA))^2)/(length(DATA)*(length(DATA)-1))
}
# T statistic
t <- qt(0.025,df=(n+m-2), lower.tail = FALSE)
term <- t*(sqrt((sample.var(X)/n)+(sample.var(Y)/m)))
# Confidence interval
print(paste0("The 95% Confidence interval for X-Y is (",
  mean(Y)-mean(X)-term,",", mean(Y)-mean(X)+term,""))
```

```
## [1] "The 95% Confidence interval for X-Y is (-0.665609515885879,0.665609515885879)"
```

```
# Bootstrap
mean.X <- c()
mean.Y <- c()
for (i in 1:B){
  # Sample of size n with replacement from data
  btstrpSmpl.X = base::sample(X,length(X), replace = TRUE)
  btstrpSmpl.Y = base::sample(Y,length(Y), replace = TRUE)
  # Mean
  mean.X[i] <- mean(btstrpSmpl.X)
  mean.Y[i] <- mean(btstrpSmpl.Y)
}
# Calculate theta
theta <- mean.X-mean.Y
# Order Theta
theta.ordered <- sort(theta)
# Where do the intervals lie
min <- (B-(0.95*B))/2
# Interval
print(paste0("The 95% Confidence interval for X-Y is (",
  theta.ordered[min],",", theta.ordered[(B-min)],"))
```

```
## [1] "The 95% Confidence interval for X-Y is (-1.54235337620315,-0.116471898835519)"
```



```
# Set parameters
B <- 200
n <- 200
mu1 <- 1
sig1 <- 1
m <- 20
mu2 <- 1
sig2 <- 2

# Generate the true samples
X <- rnorm(n,mean=mu1,sd=sig1)
Y <- rnorm(m,mean=mu2,sd=sig2)

# Theoretical Confidence interval
# Sample variance (Larsen & Marx pg 475)
sample.var <- function(DATA){
  (length(DATA)*(sum(DATA^2)) - (sum(DATA))^2)/(length(DATA)*(length(DATA)-1))
}
# T statistic
t <- qt(0.025,df=(n+m-2), lower.tail = FALSE)
term <- t*(sqrt((sample.var(X)/n)+(sample.var(Y)/m)))
# Confidence interval
print(paste0("The 95% Confidence interval for X-Y is (",
             mean(Y)-mean(X)-term,",", mean(Y)-mean(X)+term,")"))
```

```
## [1] "The 95% Confidence interval for X-Y is (-0.955542866058812,0.955542866058812)"
```

```

# Bootstrap
bootstrap.twosamp.confint <- function(DATA1,DATA2,B){
  mean.X <- c()
  mean.Y <- c()
  theta <- c()
  min <- c()
  max <- c()
  for (i in 1:B){
    # Sample of size n with replacement from data
    btstrpSmpl.X = base::sample(DATA1,length(DATA1), replace = TRUE)
    btstrpSmpl.Y = base::sample(DATA2,length(DATA2), replace = TRUE)
    # Mean
    mean.X[i] <- mean(btstrpSmpl.X)
    mean.Y[i] <- mean(btstrpSmpl.Y)
    # Calculate theta
    theta[i] <- mean.X[i]-mean.Y[i]
  }

  # Order Theta
  theta.ordered <- sort(theta)
  # Where do the intervals lie
  int <- (B-(0.95*B))/2
  # Interval
  min <- theta.ordered[int]
  max <- theta.ordered[(B-int)]
  return(c(min,max))
}

# Lets do this for B
bootstrap.conf <- bootstrap.twosamp.confint(X,Y,B)
print(paste0("The 95% Confidence interval for X-Y is (",
             bootstrap.conf[1],",", bootstrap.conf[2],")"))

```

```
## [1] "The 95% Confidence interval for X-Y is (-0.84968537443988,0.863594086427095)"
```

```

# For the plot
min.iteration <- c()
max.iteration <- c()
for (j in 1:500){
  iteration <- bootstrap.twosamp.confint(X,Y,j)
  min.iteration[j] <- iteration[1]
  max.iteration[j] <- iteration[2]
}


plot(min.iteration, type='l',
     ylim=c(-2,2),xlab="Confidence interval")
lines(max.iteration, type='l')
abline(h=(mean(Y)-mean(Y)-term), type='l', col='red')

```

```
## Warning in int_abline(a = a, b = b, h = h, v = v, untf = untf, ...): graphical  
## parameter "type" is obsolete
```

```
abline(h=(mean(Y)-mean(Y)+term), type='l', col='blue')
```

```
## Warning in int_abline(a = a, b = b, h = h, v = v, untf = untf, ...): graphical  
## parameter "type" is obsolete
```



From the plot above, it looks like the confidence interval is pretty stable once it reaches (approximately) $B=40$.